

1 L2–3 Supervised Learning

1.1 Classification and Loss

ERM: population risk $\theta^* = \arg \min_{\theta} \mathbb{E}_{x \sim p} \text{err}(f(x; \theta), y(x))$; empirical risk $\hat{\theta} = \arg \min_{\theta} \frac{1}{N} \sum \text{err}(f(x_i; \theta), y_i)$. **Binary logistic:** $p(y = 1|x) = \sigma(w^\top x + b)$,

$$\ell = -y \log p - (1 - y) \log(1 - p), \quad \frac{\partial \ell}{\partial z} = p - y.$$

Multi-class: logits $z \in \mathbb{R}^K$, $p_k = e^{z_k} / \sum_j e^{z_j}$,

$$\ell(z, y) = -\log p_y = -z_y + \log \sum_j e^{z_j}, \quad \frac{\partial \ell}{\partial z_k} = p_k - \mathbf{1}[k = y].$$

One-hot CE: $-\sum_k y_k \log p_k$.

1.2 MLP and Backprop

Layer: $h^{(0)} = x$, $z^{(\ell)} = W^{(\ell)} h^{(\ell-1)} + b^{(\ell)}$, $h^{(\ell)} = \phi(z^{(\ell)})$. Batch code convention often uses $X \in \mathbb{R}^{B \times n}$, $Z = XW + b$, $W \in \mathbb{R}^{n \times m}$. **Backprop:** define $\delta^{(\ell)} = \partial L / \partial z^{(\ell)} = \partial_{W^{(\ell)}} L = \delta^{(\ell)} (h^{(\ell-1)})^\top$, $\partial_{b^{(\ell)}} L = \delta^{(\ell)}$, $\delta^{(\ell-1)} = ((W^{(\ell)})^\top \delta^{(\ell)}) \odot \phi'(z^{(\ell-1)})$. **Activation:** sigmoid/tanh saturate; ReLU $\max(0, z)$ keeps positive-side gradient; ELU/smooth variants improve negative side. **Subgradient:** ReLU at 0 uses any value in $[0, 1]$ by convention.

1.3 CNN

Motivation: locality + weight sharing. Convolution:

$$y_{i,j,o} = \sum_{u,v,c} K_{u,v,c,o} x_{i+u,j+v,c} + b_o.$$

For input $H \times W \times C_{\text{in}}$, kernel $K_H \times K_W$, padding P , stride S : $H' = \lfloor (H + 2P - K_H)/S \rfloor + 1$, $W' = \lfloor (W + 2P - K_W)/S \rfloor + 1$. Params $K_H K_W C_{\text{in}} C_{\text{out}}$. **Scanning MLP:** same detector slides across space. **Pooling:** max/avg reduces spatial size and jitter. **Padding:** keeps boundary info; same padding for odd kernel uses $P = (K - 1)/2$.

1.4 Optimization

GD: $x_{k+1} = x_k - \eta g_k \nabla f(x_k)$. If ∇f is L -Lipschitz: $f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|y - x\|^2$, so $0 < \eta < 2/L$ decreases the upper bound. **Newton:** from quadratic Taylor, $x_{k+1} = x_k - \eta [\nabla^2 f(x_k)]^{-1} \nabla f(x_k)$. $g_k = \text{gradient on } x_k$. **Momentum:** $v_{k+1} = \beta v_k + g_k$, $x_{k+1} = x_k - \eta v_{k+1}$. **Nesterov:** compute gradient at lookahead $x_k + \beta(x_k - x_{k-1})$.

AdaGrad: $s_k = s_{k-1} + g_k^2$, $x_{k+1} = x_k - \eta g_k / (\sqrt{s_k} + \epsilon)$; may decay too fast.

RMSProp: $s_k = \rho s_{k-1} + (1 - \rho) g_k^2$.

Adam: $m_k = \beta_1 m_{k-1} + (1 - \beta_1) g_k$, $v_k = \beta_2 v_{k-1} + (1 - \beta_2) g_k^2$,

$x_{k+1} = x_k - \eta m_k / (\sqrt{\delta_k} + \epsilon)$, $\hat{m}_k = m_k / (1 - \beta_1^k)$.

1.5 Regularization and Architectures

Dropout: randomly zero units, approximates ensemble and reduces co-adaptation. **Weight decay:** add $\lambda \|\theta\|^2$. **Gradient clipping:** $g \leftarrow g \min(1, \tau / \|g\|)$.

Normalization: **GroupNorm:** batch-independent. **InstanceNorm:** batch-independent, suitable for generation tasks. **BatchNorm:** $\mu_B = \frac{1}{B} \sum_i z_i$,

$$\sigma_B^2 = \frac{1}{B} \sum_i (z_i - \mu_B)^2, \quad z_i = \gamma \frac{z_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + \beta.$$

LayerNorm: normalize hidden dimension inside one sample; better for sequence models. Batch-independent, Particularly suitable for RNN. $h^t = f\left(\frac{\alpha_t}{\sigma_t} \odot (a^t - \mu^t) + b\right)$, $\mu^t = \frac{1}{H} \sum_i a_i^t$, $\sigma_t^2 = \sqrt{\frac{1}{H} \sum_i (a_i^t - \mu^t)^2}$.

ResNet: $h_{t+1} = h_t + F(h_t)$, residual path eases optimization. **DenseNet:** concatenate previous features. **FCN:** replace fully connected layer by convolution for dense prediction.

2 L4 Energy-Based Models

2.1 Energy View

Generative modeling: learn $P(X)$ or $P(X, y) = P(y)P(X|y)$ instead of $P(y|X)$. Generative models can be used to both classify and generate images **EBM:** unnormalized score via energy

$$p_{\theta}(x) = \frac{e^{-E_{\theta}(x)}}{Z_{\theta}}, \quad Z_{\theta} = \int e^{-E_{\theta}(x)} dx.$$

Low energy means high probability; training and inference both need sampling or partition-function gradients. $\log p_{\theta}(x) = -E_{\theta}(x) - \log Z_{\theta}$, $\nabla_{\theta} \log p_{\theta}(x) = -\nabla_{\theta} E_{\theta}(x) - \frac{1}{Z_{\theta}} \nabla_{\theta} \left(\sum_{x'} e^{-E_{\theta}(x')} \right) = -\nabla_{\theta} E_{\theta}(x) + E_{x' \sim p_{\theta}} [E_{\theta}(x')]$.

We can train an energy-based model without knowing the explicit density function (or normalizing factor Z). Such as score-matching, etc.

2.2 Hopfield Network

State $y_i \in \{-1, 1\}$, local field $h_i = \sum_{j \neq i} w_{ij} y_j + b_i$. Update flips if $y_i h_i < 0$.

$$E(y) = -\sum_{i < j} w_{ij} y_i y_j - \sum_i b_i y_i = -\frac{1}{2} y^\top W y - b^\top y.$$

Asynchronous flip changes energy by $\Delta E = 2y_i h_i$; if $y_i h_i < 0$, energy decreases. Finite states imply convergence to a local minimum. **Hebbian storage:** $w_{ij} = \frac{1}{P} \sum_p y_i^p y_j^p$, $w_{ii} = 0$. Multiple patterns cause **spurious minima**; capacity is limited, roughly linear in neuron count for stable random memories. **Objective for W:** $W = \arg \min_W (\sum_{\ell \in P} E(y) - \sum_{\ell \in P} E(y'))$. **Optimization view:** desired patterns should be valleys; short-run evolution can raise undesired valleys by contrastive updates.

SGD update rule for W : $W \leftarrow W - \eta (\mathbb{E}_{y \in P[y y^\top]} - \mathbb{E}_{y' \sim p[y' y'^\top]})$, initialize y' by a random desired pattern (in P), run evolution for a few timesteps.

2.3 Boltzmann Machine

Stochastic Hopfield: energy $E(y)$ as Hopfield, $p(y) = \frac{1}{Z} e^{-E(y)} \propto e^{-E(y)}$. Conditional Gibbs update for $\{-1, 1\}$ units:

$$P(y_i = 1|y_{-i}) = \frac{1}{1 + \exp(-\sum_j w_{ij} y_j - 2b_i)} = \frac{1}{1 + e^{-2h_i}}.$$

Temperature T : $p_T(y) \propto e^{-E(y)/T}$; annealing lowers T from exploration to exploitation.

BM Training: maximize $L(W) = \sum_{y \in P} \log p_{\theta}(y)$, θ means parameters. $\nabla w_{ij} \log p_{\theta}(y) = y_i y_j - \mathbb{E}_{p_{\theta}}[y_i y_j]$, $\nabla w_{ij} L = \mathbb{E}_{y \in P}[y_i y_j] - \mathbb{E}_{y' \sim p_{\theta}}[y'_i y'_j] = \frac{1}{|P|} \sum_{y \in P} y_i y_j - \mathbb{E}_{y' \sim p_{\theta}}[y'_i y'_j]$. First term: compute average over P . Second term: sample from p_{θ} , compute average. **Sample method:** random initialize $y(0)$, cycle over i (coordinates) and sample from $P(y_i(t)|y_{-i}(t))$, after convergence select the final M states as samples. **Updating:** $w_{ij} \leftarrow w_{ij} + \eta \nabla w_{ij} L(W)$. Learning = **positive phase** data correlation minus **negative phase** model correlation. **With hidden neurons:** $y = (v, h)$

$$\nabla w_{ij} L = \frac{1}{|P|} \sum_{v \in P} \mathbb{E}_h[y_i y_j] - \mathbb{E}_{y' \sim p(v, h)}[y'_i y'_j].$$

Boltzmann for classification: $y = (v, h, c)$, c : visible, one-hot vector

2.4 RBM and Contrastive Divergence

RBM: bipartite v, h ; no visible-visible or hidden-hidden edges: $E(v, h) = -v^\top W h - b^\top v - c^\top h$. Conditionals factorize: $P(h_j = 1|v) = \sigma(c_j + \sum_i w_{ij} v_i)$, $P(v_i = 1|h) = \sigma(b_i + \sum_j w_{ij} h_j)$, $\sigma(z) = \frac{1}{1 + e^{-z}}$. Sampling: $v^{(0)} \rightarrow h^{(0)} \rightarrow v^{(1)} \rightarrow h^{(1)} \rightarrow \dots$

Maximum likelihood estimate: $\nabla w_{ij} L = \frac{1}{N_P} \sum_{v \in P} v_0 h_{ij} - \frac{1}{M} \sum v_0 h_{ij}^{\text{coj}}$.

Only 3 gibbs sampling steps needed: $\nabla w_{ij} L = \frac{1}{N_P} \sum_{v \in P} v_0 h_{ij} - v_{12} h_{ij}$.

CD-k: start at data $v^{(0)}$, run k Gibbs steps, update $W \propto v^{(0)} h^{(0)\top} - v^{(k)} h^{(k)\top}$.

2.5 Sampling

Inverse CDF: sample $u \sim U[0, 1]$, choose category by cumulative mass. **Box-Muller:** Gaussian from uniform:

$$z_1 = \sqrt{-2 \log u_1} \cos(2\pi u_2), \quad z_2 = \sqrt{-2 \log u_1} \sin(2\pi u_2).$$

Importance sampling: $\mathbb{E}_p[f(x)] = \mathbb{E}_q[f(x)p(x)/q(x)]$. Good q must cover p 's high-mass region. **MCMC:** Detailed balance $\pi(x)T(x \rightarrow x') = \pi(x')T(x' \rightarrow x)$ implies stationarity of π . **Ergodic:** (intuitively) you can visit any desired state with positive probability from any state, ensure a unique stationary distribution. **Metropolis Hastings Algo.:** Construct Markov chain with stationary π (desired distribution). Current x , draw $x' \sim q(x'|x)$ (called proposal), transit to x' with prob. $\alpha = \min\left(1, \frac{\pi(x')q(x|x')}{\pi(x)q(x'|x)}\right)$, otherwise stays at x . Choice of q : random proposal $q(x'|x)$ =noise (e.g. Gaussian, then it relies on $|x - x'|$), then q terms in α vanishes.

For EBM, partition Z cancels. **Gibbs:** sample a single coordinate per step. Acceptance ratio always 1. **SGLD:** $x_{t+1} = x_t - \frac{\eta}{2} \nabla_x E_{\theta}(x_t) + \sqrt{\eta} \xi_t$.

3 L5 Generative Adversarial Networks

3.1 Implicit Generative Model

Instead of density $p_{\theta}(x)$, learn sampler $x = G_{\theta}(z)$, $z \sim p_z$. Likelihood-free learning. Need differentiable distance between generated and real distributions. GAN learns this distance with a discriminator. **GAN:** generator $G_{\theta}(z)$, $z \sim p(z) = N(0, I)$; discriminator $D_{\phi}(x)$, classify the

data is from p_{data} or G_{θ} . Objective: $L = \min_{\theta} \max_{\phi} V(D_{\phi}, G_{\theta})$.

$$V(D_{\phi}, G_{\theta}) = \mathbb{E}_{x \sim p_{\text{data}}} \log D_{\phi}(x) + \mathbb{E}_{z \sim p_z} \log(1 - D_{\phi}(G_{\theta}(z))).$$

Dataset: $\{(x, 1) : x \sim p_{\text{data}}\} \cup \{(\tilde{x}, 0) : \tilde{x} \sim G_{\theta}(z)\}$.

Train D: $L(\phi) = \mathbb{E}_{x \sim p_{\text{data}}} \log D_{\phi}(x) + \mathbb{E}_{z \sim p_z} \log(1 - D_{\phi}(G(z)))$.

Train G: $L(\theta) = \mathbb{E}_{z \sim p(z)} [\log D_{\phi}(G_{\theta}(z))]$.

For fixed G , the optimal D is $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + P_G(x)}$. Substitute:

$$V(D^*, G) = -\log 4 + 2\text{JSD}(p_{\text{data}} \| p_G), \quad \text{JSD}(\|p\|q) = \text{KL}(p \| \frac{p+q}{2}) + \text{KL}(q \| \frac{p+q}{2}), \geq 0, \quad \sqrt{\text{JSD}} \text{ satisfies triangular inequality.}$$

$$2\text{KL}(p\|q) \geq (\int |p(x) - q(x)| dx)^2.$$

Non-saturating generator loss $-\mathbb{E}_z \log D(G(z))$ gives stronger gradients.

3.2 Evaluation and Failure Modes

Mode collapse: generator covers few modes (converges to a few samples). Intrinsic issue of JSD. If we optimize q_{θ} w.r.t. a multi-modal distribution p using KL-divergence $\text{KL}(q_{\theta} \| p)$, then **Training instability:** moving target game; discriminator too strong $\rightarrow$ weak generator gradient. **IS:** rewards confident/diverse ImageNet predictions, but blind to intra-class artifacts. **FID:** Compute μ_p, Σ_p and μ_G, Σ_G for $p_{\text{data}}, G(z)$ using inception v3 pool3 layer (2048-d), Compute Wasserstein distance between two Gaussians. compare Gaussian feature statistics: $\|\mu_p - \mu_g\|^2 + \text{tr}(\Sigma_p + \Sigma_g - 2(\Sigma_p \Sigma_g)^{1/2})$. Lower FID score means better generators.

3.3 WGAN Family

JSD can be flat when supports do not overlap. Wasserstein-1: $W(p, q) = \sup \|f\|_{L \leq 1} \mathbb{E}_{x \sim p} f(x) - \mathbb{E}_{x \sim q} f(x)$.

WGAN: critic f_{ϕ} , no D . WGAN objective: subject to $\|f_{\phi}(x)\|_L \leq 1$, $L(\phi) = \mathbb{E}_{x \sim p_{\text{data}}} [f_{\phi}(x)] - \mathbb{E}_{x \sim P_G} [f_{\phi}(x)]$, $L(\theta) = \mathbb{E}_{x \sim G_{\theta}(z)} [f_{\phi}(x)]$.

Lipschitz constraint: weight clipping, spectral norm, or **WGAN-GP:** $\tilde{x} \leftarrow (1 - \epsilon)x + \epsilon \tilde{x}$, $\tilde{\epsilon} \sim \text{Unif}(0, 1)$, $L(\phi) = \mathbb{E}_{x \sim p_{\text{data}}} [f_{\phi}(x)] - \mathbb{E}_{\tilde{x} \sim P_G} [f_{\phi}(\tilde{x})] + \lambda \mathbb{E}_{\tilde{x}} (\|\nabla_{\tilde{x}} f_{\phi}(\tilde{x})\|_2 - 1)^2$.

DCGAN: conv/deconv, batchnorm, stable architecture. **BigGAN:** scale + class conditioning + truncation trick. **GigaGAN:** text-to-image GAN scaling. **R3GAN:** regularized robust objective.

3.4 Variants and Applications

Conditional GAN: $G(z, c)$, $D(x, c)$. **Semi-supervised GAN:** discriminator has $K + 1$ classes, K real labels plus fake. **BIGAN/ALI:** learn encoder $E(x)$ with discriminator over (x, z) pairs. Idea: $(z, G(z))$ and $(E(x), x)$ from same distribution. gives representation learning. **Pix2Pix:** paired image-to-image translation with conditional GAN + reconstruction loss. **CycleGAN:** unpaired translation using cycle consistency $G : X \rightarrow Y$, $F : Y \rightarrow X$, $\|F(G(x)) - x\| + \|G(F(y)) - y\|$. **DragGAN:** interactive image manipulation by optimizing latent/code with handle-point constraints.

4 L6 Flow and VAE

4.1 Normalizing Flow

Invertible differentiable generator $x = G_{\theta}(z)$, transform a simple distribution to the data distribution. $\log p_X(x) = \log p_Z(G^{-1}(x)) + \log \left| \det \frac{\partial G^{-1}}{\partial x} \right|$.

Composition: $z_K = G(z_0) = f_K \circ \dots \circ f_1(z_0)$, $\log p_X(z_K) = \log p_Z(z_0) - \sum_k \log |\det \frac{\partial f_k}{\partial z_k}|$. Flow design needs invertible layer, efficient inverse, efficient compute determinant of Jacobi.

Coupling: split $x = (x_a, x_b)$, to (y_a, y_b) . $y_a = x_a$, $y_b = x_b \odot e^{s(x_a) + t(x_b)}$, $\log |\det J| = \sum s(x_a)$. s, t : learnable. Sometimes needs premutation among coordinates or linear transformation.

Glow: invertible 1×1 convolution mixes channels; LU decomposition makes determinant cheap: $\log |\det W| = \sum_i \log |u_{ii}|$. **Dequantization:** add uniform noise to discrete pixels so continuous density is well-defined. **Continuous NF / Neural ODE:** $\frac{dz(t)}{dt} = f_{\theta}(z(t), t)$, $\log p(z(t_1)) = \log p(z(t_0)) - \int_{t_0}^{t_1} \text{tr} \left(\frac{\partial f_{\theta}}{\partial z(t)} \right) dt$.

Prove a Normalizing Flow: 1. invertible; 2. easily-tractable Jacobi matrix, **4.2 EM and Variational Inference**

Latent-variable likelihood $\log p_{\theta}(x) = \log \sum_z p_{\theta}(x, z)$. Jensen:

$$\log p_{\theta}(x) = \log \mathbb{E}_q \frac{p_{\theta}(x, z)}{q(z)} \geq \mathbb{E}_q \log p_{\theta}(x, z) + H(q) = \mathcal{L}(q, \theta).$$

Gap: $\log p_{\theta}(x) - \mathcal{L}(q, \theta) = \text{KL}(q(z) \| p_{\theta}(z|x))$. EM: E-step $q(z) = p_{\text{old}}(z|x)$, M-step maximize $\mathbb{E}_q \log p_{\theta}(x, z)$.

4.3 VAE

Encoder $q_{\phi}(z|x)$, decoder $p_{\theta}(x|z)$, prior $p(z)$.

ELBO: $\log p_{\theta}(x) \geq \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] - \text{KL}(q_{\phi}(z|x) \| p(z))$. In practice, $q_{\phi}(z|x) = \mathcal{N}(\mu_{\phi}(x), \sigma_{\phi}(x))$, $p_{\theta}(x|z) = \mathcal{N}(\mu_{\theta}(z), \beta I)$. Reparameterization: $z = \mu_{\phi}(x) + \sigma_{\phi}(x) \odot \epsilon$, $\epsilon \sim \mathcal{N}(0, I)$. KL of Gaussian can be computed, here $\text{KL}(\mathcal{N}(\mu, \sigma^2 I) \| \mathcal{N}(0, I)) = \frac{1}{2} \sum_i (\mu_i^2 + \sigma_i^2 - \log \sigma_i^2 - 1)$.

Terms: reconstruction = data fidelity; KL = latent regularization. In β -VAE,

large β enforces latent variables to be less correlated with each other:

$$\mathcal{L}_{\beta\text{-VAE}} = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)] - \beta \text{KL}(q_{\phi}(z|x) \| p(z)).$$

PixelVAE/NVAE: hierarchical latent variables improve global coherence and likelihood.

5 L7 Diffusion and Flow Matching

5.1 DDPM

Forward: $q(x_t|x_{t-1}) = \mathcal{N}(\sqrt{\alpha_t}x_{t-1}, \beta_t I)$, $\beta_t = 1 - \alpha_t$, $\hat{\alpha}_t = \prod_{s=1}^t \alpha_s$. Closed form: $q(x_t|x_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$, $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$. Reverse: use $p_{\theta}(x_{t-1}|x_t)$ to estimate the prob, $p_{\theta}(0 : T) = p(x_T) \prod_t p_{\theta}(x_{t-1}|x_t)$. $p(x_T)$ is gaussian prior.

ELBO $\log p_{\theta}(x_0) \geq \mathbb{E}_{x_1: T \sim q(\cdot|x_0)} [\log \frac{p_{\theta}(x_0:T)}{q(x_1:T|x_0)}]$, use markovian and use tractable $q(x_{t-1}|x_t, x_0)$: $L = \mathbb{E}_q [\log \frac{p(x_T)}{q(x_T|x_0)}] + \mathbb{E}_q [\log p_{\theta}(x_0|x_1)] + \sum_{t=2}^T \text{KL}(q(x_{t-1}|x_t, x_0) \| p_{\theta}(x_{t-1}|x_t))$ (core of DDPM).

Posterior $q(x_{t-1}|x_t, x_0) = \frac{1}{q(x_t|x_0)} q(x_t|x_{t-1}, x_0) q(x_{t-1}|x_0)$, which is a

Gaussian $N(\bar{\mu}_t, \bar{\beta}_t I)$. Let $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$, $\bar{\mu}_t = \frac{\sqrt{\bar{\alpha}_t}x_0}{1 - \bar{\alpha}_t}x_0 + \frac{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_t)}{1 - \bar{\alpha}_t}x_t = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(x_t - \frac{1 - \bar{\alpha}_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon \right)$, $\bar{\beta}_t = \frac{1 - \bar{\alpha}_t - 1}{1 - \bar{\alpha}_t} \beta_t$.

Predict noise $\epsilon_{\theta}(x_t, t)$: $p_{\theta}(x_{t-1}|x_t) = \mathcal{N}(\mu_{\theta}(x_t, t), \sigma_{\theta}^2 I)$. $\mu_{\theta}(x_t, t) = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(x_t - \frac{1 - \bar{\alpha}_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right)$. Training: sample x_0 from data, sample timestep t , add noise ϵ , predict $\epsilon_{\theta}(x_t, t)$. Loss:

$$L = \mathbb{E}_{x_0 \sim p_{\text{data}}, t \sim [1:T], \epsilon \sim \mathcal{N}(0, I)} \|\epsilon_{\theta}(x_t, t) - \epsilon\|_2^2, \quad x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon.$$

$$\text{Sampling: } x_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(x_t, t) \right) + \sigma_t z.$$

5.2 Score-Based Models

Score: learn score $\nabla_x \log p(x)$. **Langevin dynamics:** use only score. $x_{t+1} = x_t + \eta \nabla_x \log p(x_t) + \sqrt{2\eta} \epsilon_t$, $\epsilon_t \sim \mathcal{N}(0, I)$. continuous form: $dx_t = \nabla_x \log p(x_t) dt + \sqrt{2dwt_t} dw_t$ ($N(0, d_t)$). **NCSN:** noise conditional score network. $x_i = x_0 + \sigma_i \epsilon$, learn $S_{\theta}(x_i, \sigma_i)$ to estimate $\nabla_x \log p(x_i)$.

$x_t = \alpha_t x_0 + \sigma_t \epsilon$, score $S_{\theta}(x_t, t)$? MSE loss $L = \mathbb{E}_{x_t} \|S_{\theta}(x_t, t) - \nabla_{x_t} \log p(x_t)\|^2$. $\mathbb{E}_{x_t} [S_{\theta}(x_t, t)^\top \nabla_{x_t} \log p(x_t)] = \mathbb{E}_{x_t} [\nabla \cdot S_{\theta}(x_t, t)]$, using $\nabla \cdot (sp) = (\nabla \cdot s)p + s^\top \cdot \nabla p$. **Denoisng score matching:** $\nabla_{x_t} \log p(x_t) = \mathbb{E}_{x_0 \sim p(\cdot|x_t)} [\nabla_{x_t} \log p(x_t|x_0)]$

$\mathbb{E}_{x_t} \|S_{\theta}(x_t, t) - \nabla_{x_t} \log p(x_t|x_0)\|^2$, $L_{\text{DSM}} = \mathbb{E}_{x_0 \sim p_{\text{data}}, t \sim [0, 1], x_t \sim q(x_t|x_0)} \|S_{\theta}(x_t, t) - \nabla_{x_t} \log q(x_t|x_0)\|^2$, $q(x_t|x_0)$ is tractable. If $x_t = x_0 + \sigma_t \epsilon$, $\nabla_{x_t} \log q(x_t|x_0) = \nabla_{x_t} (x_t - \frac{(x_t - \alpha_t x_0) \sigma_t^2}{2\sigma_t^2}) = -\frac{x_t - \alpha_t x_0}{\sigma_t^2} = -\frac{\epsilon}{\sigma_t}$. Let $S_{\theta}(x_t, t) = -\frac{\epsilon_{\theta}(x_t, t)}{\sigma_t}$, loss same as DDPM.

SDE: $dx = f(x, t)dt + g(t)dW_t$, reverse: $dx = [f(x, t) - g(t)^\top \nabla_x \log p_t(x)]dt + g(t)d\bar{W}_t$. A equivalent ODE (same distribution when same t): Probability-flow ODE: $dx = [f(x, t) - \frac{1}{2} g(t)^2 \nabla_x \log p_t(x)]dt$. Likelihood along ODE: $\log p_0(x_0) =$

$$\log p_T(x_T) + \int_0^T \text{tr} \left(\frac{\partial f_t}{\partial x_t} \right) dt.$$

Inference: use learned score to solve SDE/ODE using a solver. Hence #sample-steps in training and inference can be different. DDIM: sample according to schedule

5.3 DPM-Solver and Applications

DPM solver 1 order DPM = DDIM. DPM-Solver integrates diffusion ODE with high

6.1 Sequence Data and Tokenization

Sequence data includes language, audio, action; images/DNA/protein can be serialized. Autoregressive factorization: $p(x_{1:T}) = \prod_{t=1}^T p(x_t|x_{<t})$. **Tokenization**: text $\rightarrow$ basic units $\rightarrow$ token IDs. Byte-Pair Encoding (BPE) repeatedly merges frequent adjacent symbols; balances vocabulary size and sequence length. **Embedding**: token ID $i \rightarrow e_i = W_E[i]$; positional info is needed because token embeddings alone are order-free.

6.2 Autoregressive CNN and Pixel Models

WaveNet: causal/dilated temporal convolution expands receptive field without recurrence. **PixelCNN**: image likelihood $p(x) = \prod_i p(x_i|x_{<i})$. under raster order; masked convolution prevents access to future pixels. Gated PixelCNN uses gates to improve expressiveness.

6.3 RNN and LSTM

RNN: $h_t = \tanh(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$, $y_t = W_{hy}h_t + b_y$. BPTT multiplies many Jacobians; singular values < 1 vanish, > 1 explode. **LSTM**: $f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f)$, $i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i)$, $\tilde{c}_t = \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$, $c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$, $o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o)$, $h_t = o_t \odot \tanh(c_t)$. Additive cell path gives controllable memory; forget gate decides retention.

6.4 Attention and Transformer

For token t : $q_t = W_Qx_t$, $k_j = W_Kx_j$, $v_j = W_Vx_j$, $\alpha_{tj} = \frac{\exp(q_t^\top k_j / \sqrt{d_k})}{\sum_i \exp(q_t^\top k_i / \sqrt{d_k})}$, $y_t = \sum_j \alpha_{tj} v_j$. Matrix form: $\text{Attn}(Q, K, V) = \text{softmax}(QK^\top / \sqrt{d_k})V$. **Multi-head**: parallel attention subspaces, concatenate and project. **Causal mask**: block future tokens. **Transformer block**: attention + MLP + residual + normalization. Complexity $O(T^2d)$; FlashAttention reduces memory IO, Mamba/SSM explores linear-time sequence mixing.

6.5 Objectives and Case Studies

Objectives: autoregressive LM, masked/denoising, contrastive predictive. **Recipe**: pretraining, fine-tuning, low-rank tuning. **GPT**: decoder-only causal Transformer; next-token pretraining gives in-context learning with scale. **ViT**: image patches as tokens. **Swin**: shifted local windows for hierarchical vision. **CLIP**: contrastive image-text alignment,

$$L = -\frac{1}{2} \left[\frac{1}{N} \sum_i \log \frac{e^{s_{ii}/\tau}}{\sum_j e^{s_{ij}/\tau}} + \frac{1}{N} \sum_i \log \frac{e^{s_{ii}/\tau}}{\sum_j e^{s_{ji}/\tau}} \right].$$

SigLIP: pairwise sigmoid loss avoids global softmax. **Whisper**: audio Transformer trained on large-scale speech. **MAR**: masked autoregressive modeling for images.

7 L10 Modern Large Models

7.1 Pretraining and Scaling

Autoregressive LM: maximize $\sum_t \log p_\theta(x_t|x_{<t})$. Base model learns continuation; chat model learns instruction-following behavior through post-training. **Scaling laws**: loss decreases smoothly with model size N , data D , compute C , often power-law: $L(N) \approx L_\infty + aN^{-\alpha}$, $L(D) \approx L_\infty + bD^{-\beta}$. Compute-optimal training balances parameters and tokens; not simply largest model on fixed data. Larger models reach the same loss with fewer samples / fewer optimization steps. **Prompting**: zero-shot/few-shot describes task in context; no parameter update. **Instruction tuning**: supervised fine-tuning on instruction-response mixtures.

7.2 Post Training

Using human feedback. **SFT**: next-token imitation on human demonstrations. It cannot express preference among multiple valid answers. **RLHF**: human preference. Can reward uncertainty awareness, e.g. "I don't know" when question is unanswerable. Negative samples provide learning signal unavailable in pure SFT. Helps eliminating hallucination. **Reward model**: train from pairwise rankings:

$$L_{\text{RM}} = -\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l)).$$

Policy optimization with KL:

$$\max_{\pi} \mathbb{E}_{y \sim \pi(\cdot|x)} [r_\phi(x, y)] - \beta \text{KL}(\pi(\cdot|x) \parallel \pi_{\text{ref}}(\cdot|x)).$$

KL prevents policy from drifting too far from SFT/reference.

7.3 RAG and Reasoning

RAG: retrieve documents d then generate:

$$p(y|x) \approx \sum_{d \in \mathcal{R}(x)} p_\theta(y|x, d)p(d|x).$$

Retrieval supplies missing knowledge but does not automatically solve multi-step reasoning. **Inference-time scaling**: spend more test-time compute on search, verification, self-consistency, or long reasoning trajectories. For reasoning model:

$$\max_{\pi} \mathbb{E}_{\tau \sim \pi(\cdot|x)} [R(\text{answer}(\tau), y)] - \beta \text{KL}(\pi \parallel \pi_{\text{ref}}).$$

Answer verifier gives sparse reward; RL can optimize thinking trajectory even without unique chain-of-thought labels.

7.4 Multimodal Models

Vision-language few-shot: convert image/video tasks to text generation conditioned on visual tokens. **Flamingo**: frozen vision encoder + frozen LM with cross-attention layers and Perceiver resampler for interleaved image/video-text. **LLaVA**: CLIP vision encoder + projection to Vicuna/LLM; visual instruction tuning teaches chat-style responses. **Unified multimodal model**: one decoder handles understanding, generation, editing, reasoning; may use separate experts/tokenizers for text vs image. **BAGEL-style idea**: decoder-only unified model with modality-specific experts; understanding often emerges before reasoning/editing, which need larger scale.

8 Cross-Lecture Map

Density models: EBM defines $p \propto e^{-E}$ but sampling is hard; flow gives exact likelihood via invertible map; VAE gives ELBO; diffusion learns reverse denoising, flow matching learns vector field directly. **Implicit models**: GAN learns sampler without tractable likelihood. **Sequence models**: Transformer turns all modalities into token sequences. **Large models**: pretraining builds capability, post-training aligns behavior, inference-time compute and multimodal grounding extend usability.